

Data Storage

Philip J. Cwynar

University of Pittsburgh
School of Information Sciences
pcwynar@pitt.edu

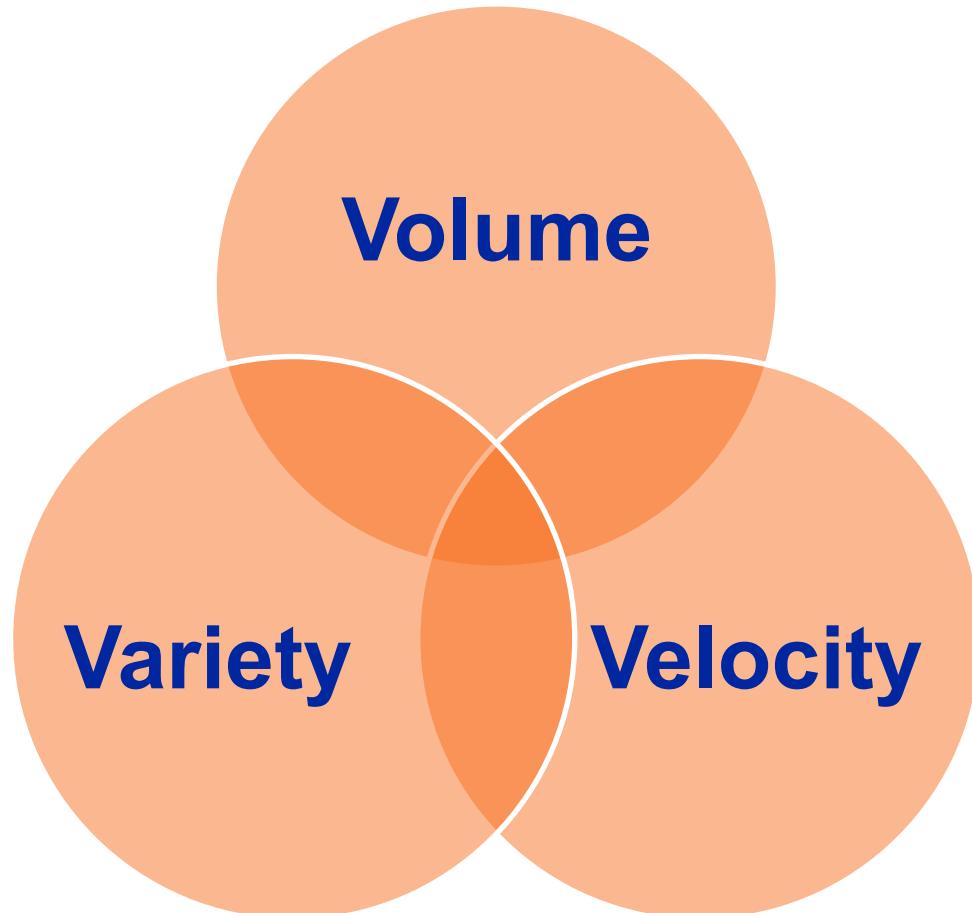
Outline

- Collecting data
- Relational Databases
- NoSQL Databases
 - Key-value databases
 - Document databases
 - Column-family stores
 - Graph databases
- Beyond NoSQL Databases

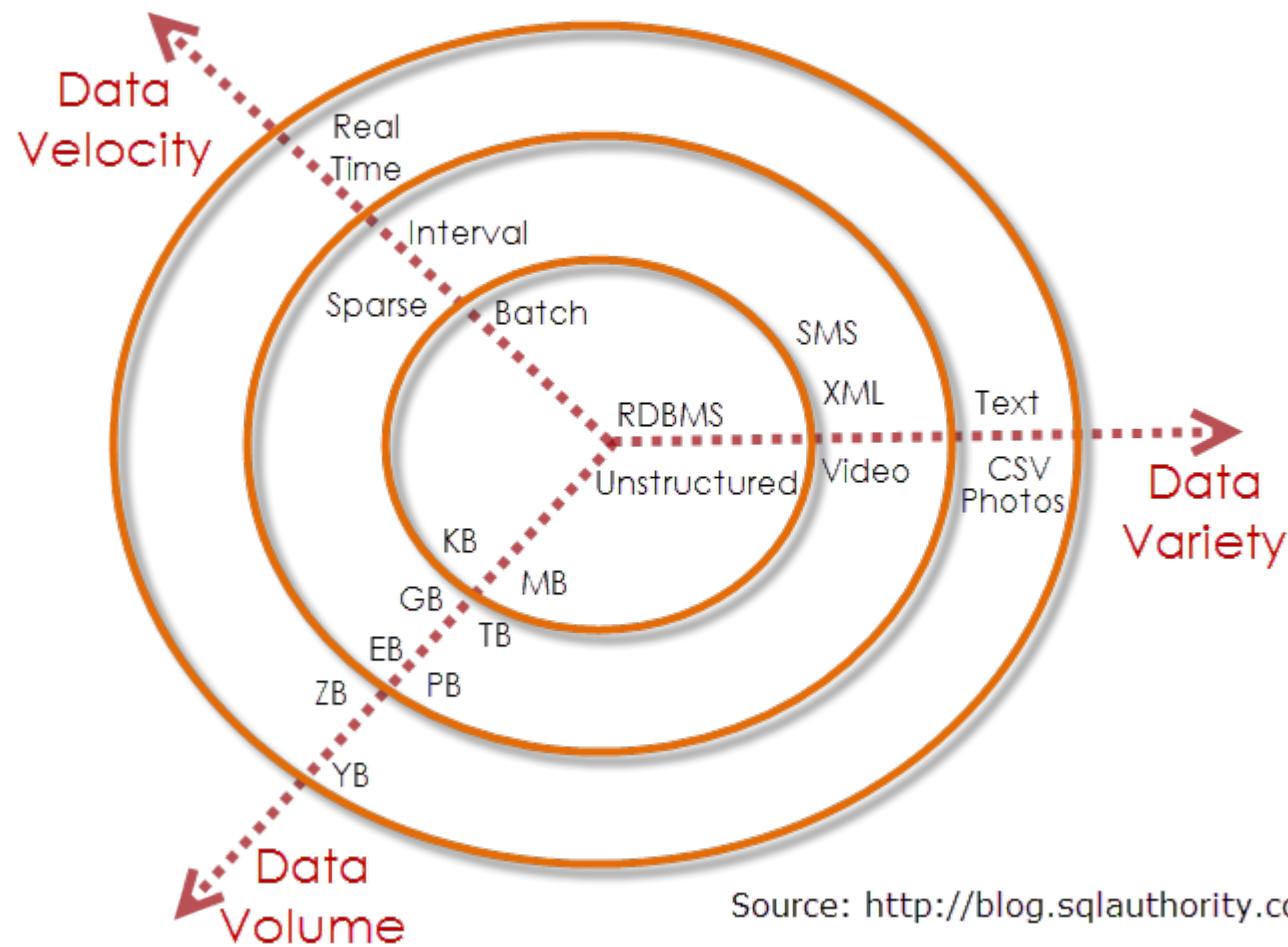
Collecting Data

- Collecting data
- Relational Databases
- NoSQL Databases
- Beyond NoSQL Databases

Big Data's 3V



3Vs of Big Data



The quantity of data collected

- The New York Stock Exchange generates about one terabyte of new trade data per day.
- Facebook hosts approximately 10 billion photos, taking up one petabyte of storage.
- Ancestry.com, the genealogy site, stores around 2.5 petabytes (10^{15}) of data.
- The Internet Archive stores around 2 petabytes (10^{15}) of data, and is growing at a rate of 20 terabytes (10^{12}) per month.
- The Large Hadron Collider near Geneva, Switzerland, will produce about 15 petabytes (10^{15}) of data per year.

Reference: Tom White, Hadoop: The Definitive Guide, Third Edition, 2012

Structured Data

I suspect the term SD came from the name for a common language used to access DBs, called Structured Query Language, or [SQL](#).

SQL provides a well-defined way for applications to manage data in a DB.

The little snippet of SQL code here illustrates how an application could retrieve all rows from a table called Book where the Price is greater than 100 and request that the result be sorted in ascending order by title.

```
SELECT *
  FROM Book
 WHERE price > 100
 ORDER BY title;
```

Storing Data: Relational Databases

Relational databases: Definition

- **Relational database: A set of relations**
- **Relation: Made up of 2 parts:**
 - **Schema: specifies name of relation plus name and type of each column**
`Customer(id:int, name:string, gender:string, email:string)`
 - **Instance: a table, with rows and columns**

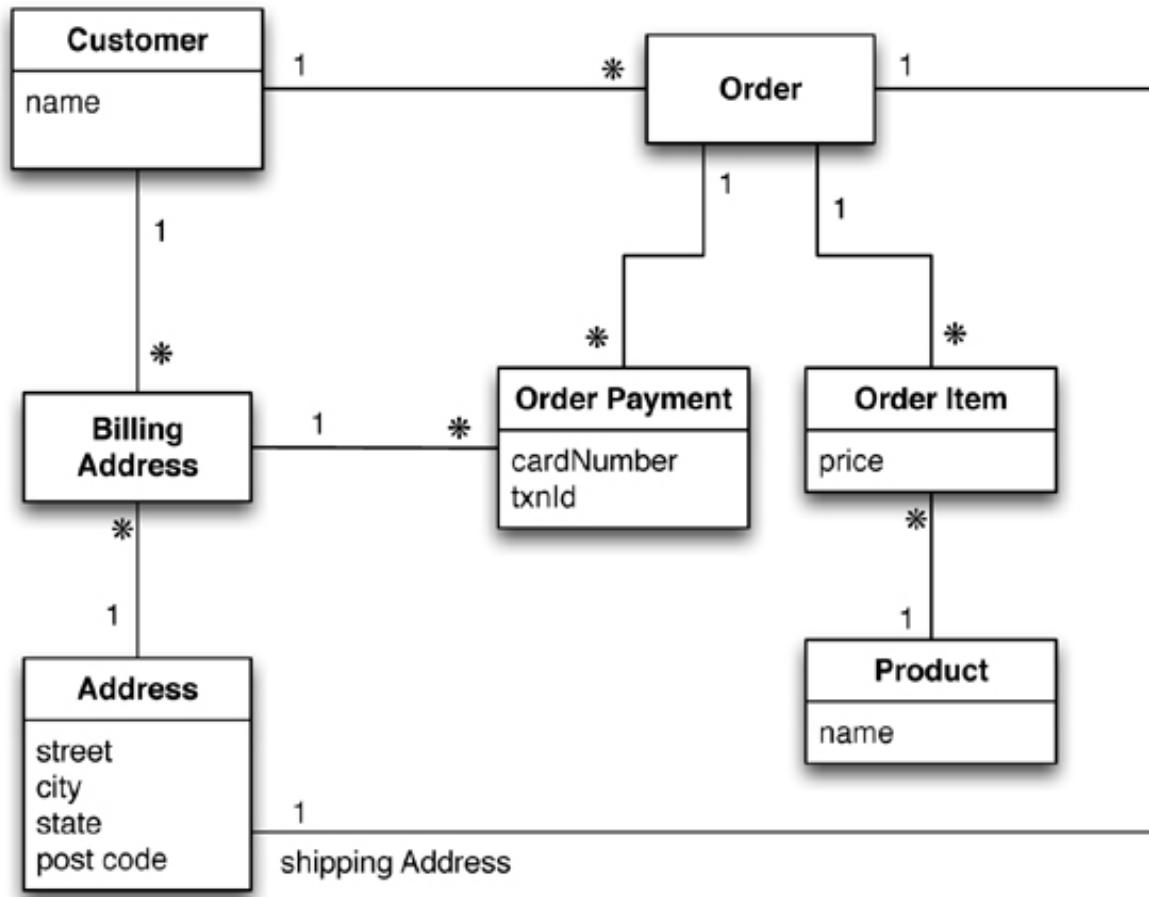
id	name	gender	email
1	Philip	Male	pcwynar@pitt.edu
2	Chirayu	Male	chw@pitt.edu

- **Can think of a relation as a set of rows (tuples)**

Reference for this section: NoSQL distilled: A brief guide to the emerging world of polyglot persistence / Pramod J Sadalage, Martin Fowler.

Relational databases: Data Model

Entity Relationship Diagram



Data model oriented around a relational database (using UML notation)

Relational databases: Data Model

Customer	
Id	Name
1	Martin

Orders		
Id	CustomerId	ShippingAddressId
99	1	77

Product	
Id	Name
27	NoSQL Distilled

BillingAddress		
Id	CustomerId	AddressId
55	1	77

OrderItem			
Id	OrderId	ProductId	Price
100	99	27	32.45

Address	
Id	City
77	Chicago

OrderPayment				
Id	OrderId	CardNumber	BillingAddressId	txnId
33	99	1000-1000	55	abelif879rft

Typical data using RDBMS data model

Relational databases: Integrity constraints (ICs)

- IC: condition that must be true for any instance of the database, e.g., domain constraints
 - ICs are specified when schema is defined
 - ICs are checked when relations are modified
- A legal instance of a relation is one that satisfies all specified ICs.
 - DBMS should not allow illegal instances

Relational databases: Primary key constraints

- A set of fields is a key for a relation if:
 1. No two distinct tuples can have same value in all key fields, and
 2. This is not true for any subset of the key
 - » If part 2 is false, then it is super key.
- K is a candidate key if K is minimal.
- Among all candidate keys we must select one that becomes the Primary Key

Relational databases: Foreign key constraints

- **Foreign key:** Set of fields in one relation that is used to refer to a tuple in another relation (must correspond to a primary key of the second relation)
- If all foreign key constraints are enforced, referential integrity is achieved

Customer	
Id	Name
1	Martin

Orders		
Id	CustomerId	ShippingAddressId
99	1	77

Relational databases: Normalization

- **1st Normal Form:** make tables flat (all elements atomic)
- **2nd Normal Form:** Every non-prime attribute depends on a candidate key or another non-prime attribute
- **3rd Normal Form:** some redundancy but dependency preserving
- **BCNF (Boyce Codd Normal Form):** no redundancy but not dependency preserving
- ...
- **Redundancy might lead to update anomalies**
- **Note:** we are going to break normalization (denormalize) later

Relational databases: Query language

- A major strength of the relational model:
 - Supports simple, powerful querying of data (**SQL**)
 - » **DDL** (Data Description Language): create, drop, alter
 - » **DML** (Data Manipulation Language): insert, update, delete, select
 - » **DCL** (Data Control Language): grant, revoke
 - » **TCL** (Transaction Control Language): commit, rollback
- Queries can be written intuitively, and the DBMS is responsible for efficient evaluation
 - The key: precise semantics for relational queries
 - Allows the optimizer to extensively re-order operations, and still ensure that the answer does not change.

Relational databases: Example

Import ODBC data

Tables Views SQL Query

Genres MovieGenres Movies

ID
 Year
 Title
 r1
 r2
 r3
 r4
 r5
 p1
 p2
 p3
 p4
 p5

Make Query

ID	Year	Title	r1	r2	r3	r4	r5	p1
1	2003	Dinosaur Planet	28	31	136	207	145	0.05118
2	2004	Isle of Man TT 2004 Review	16	11	35	42	41	0.1103
3	1997	Character	73	153	554	875	357	0.03628
4	1994	Paula Abdul's Get Up & Dance	30	34	40	19	19	0.2112
5	2004	The Rise and Fall of ECW	118	51	140	327	504	0.1035
6	1997	Sick	173	136	295	262	153	0.1697
7	1992	8 Man	32	30	21	7	3	0.3440
8	2004	What the #\\$! Do We Know!?	2089	2505	3766	3587	2963	0.140107
9	1991	Class of Nuke 'Em High 2	21	25	28	11	10	0.221053
10	2001	Fighter	25	34	93	65	32	0.100402
11	1999	Full Frame: Documentary Shorts	20	38	75	46	19	0.10101
12	1947	My Favorite Brunette	30	56	209	158	93	0.0549451
13	2003	Lord of the Rings: The Return of the King	0	1	9	35	80	0.0
14	1982	Nature: Antarctica	17	17	43	28	13	0.144068
15	1988	Neil Diamond: Greatest Hits Live	30	35	88	96	41	0.103448
16	1996	Screamers	158	511	1146	675	209	0.0585402
17	2005	7 Seconds	669	1596	3013	1414	416	0.0941193

Movies

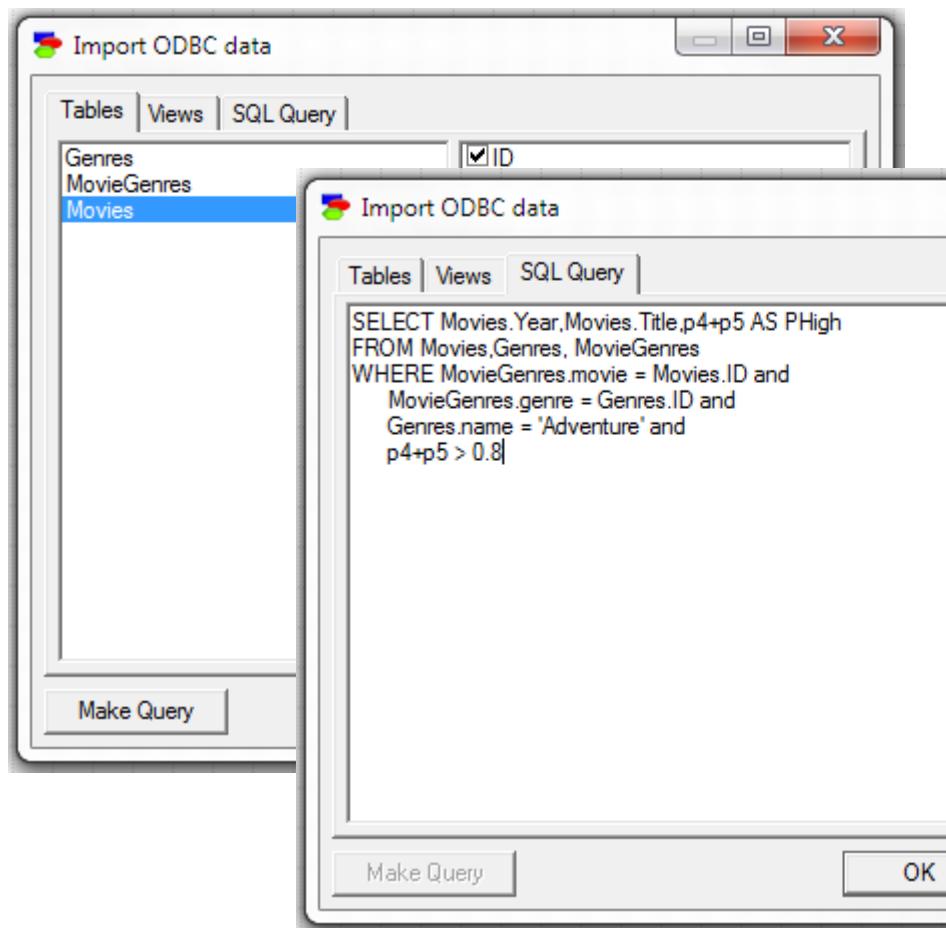
Genres

ID	Name
12	Action
1	Adult
16	Adventure
10	Animation
22	Biography
6	Comedy
13	Crime
7	Documentary
2	Drama
11	Family
14	Fantasy
24	Film-Noir
25	Game-Show
8	History

MovieGenres

movie	genre			
13519	1			
12056	1			
14249	2			
6745	3			
6745	4			
6745	5			
12540	6			
12540	4			
15399	6			
5183	7			
15760	7			
15760	3			
8769	7			
8769	8			
8769	9			
2945	6			
2945	4			
3921	10			
3921	6			
3921	11			
13240	6			
11625	6			
10198	6			
10198	11			
10198	11			
0.248629	0.378428	0.265082		
0.241379	0.289655	0.282759		
0.275348	0.434891	0.177435		
0.28169	0.133803	0.133803		
0.122807	0.286842	0.442105		
0.2895	0.257115	0.150147		
0.225806	0.0752688	0.0322581		
0.252582	0.240577	0.198726		
0.263158	0.294737	0.115789	0.105263	
0.373494	0.261044	0.128514		
0.378788	0.232323	0.0959596		
0.382784	0.289377	0.17033		
0.008	0.072	0.28	0.64	
0.144068	0.364407	0.237288	0.110169	
0.12069	0.303448	0.331034	0.141379	
0.189329	0.424602	0.250093	0.0774361	
0.224536	0.423889	0.198931	0.0585256	

Relational databases: Query Example



	Year	Title	PHigh
▶	1959	North by Northwest	0.809471
	1993	Star Trek: The Next Generation: Season 7	0.836429
	1954	Seven Samurai	0.800669
	2003	Read Or Die	0.827824
	2001	Alias: Season 1	0.846131
	2004	Samurai Champloo	0.8825
	1991	Star Trek: The Next Generation: Season 5	0.834596
	2004	Case Closed: Season 5	0.866141
	2004	Teen Titans: Season 2	0.863637
	2004	Lost: Season 1	0.941647
	1992	Star Trek: The Next Generation: Season 6	0.846714
	2000	Farscape: Season 2	0.81321
	2005	Batman Begins	0.829358
	2004	Farscape: The Peacekeeper Wars	0.831399
	2002	Farscape: Season 4	0.840428
	2002	Alias: Season 2	0.857559
	1990	Star Trek: The Next Generation: Season 4	0.84949
	2003	Alias: Season 3	0.849061
	2002	Smallville: Season 2	0.80966
	1987	The Princess Bride	0.832354
	2004	Battlestar Galactica: Season 1	0.93818
	1995	Ninja Scroll	0.804775
	1981	Raiders of the Lost Ark	0.908969
	2003	Smallville: Season 3	0.813642
	2004	The Incredibles	0.842713
	1963	The Great Escape	0.816221
	2004	Harry Potter and the Prisoner of Azkaban	0.807584
	2001	Farscape: Season 3	0.844224
	1995	Toy Story	0.858613
	2000	Gladiator	0.805826
	1999	Toy Story 2	0.809113
	1988	Star Trek: The Next Generation: Season 2	0.824505
	2004	Smallville: Season 4	0.874619
	1989	Indiana Jones and the Last Crusade	0.852522

Relational databases: ACID transactions

- Transaction is a unit of work performed within a database management system against a database
- **Atomicity** – all or nothing
- **Consistency** – bring database from one valid state to another
- **Isolation** - the intermediate state of a transaction is invisible to other transactions
- **Durability** - after a transaction successfully completes, changes to data persist and are not undone, even in the event of a system failure.

Structured Data Warehousing Approaches

Bill Inmon – “Father of Data Warehousing”
----Enterprise Data Warehouse

Ralph Kimball – “Father of Business Intelligence”
----Data Mart / Star Schema

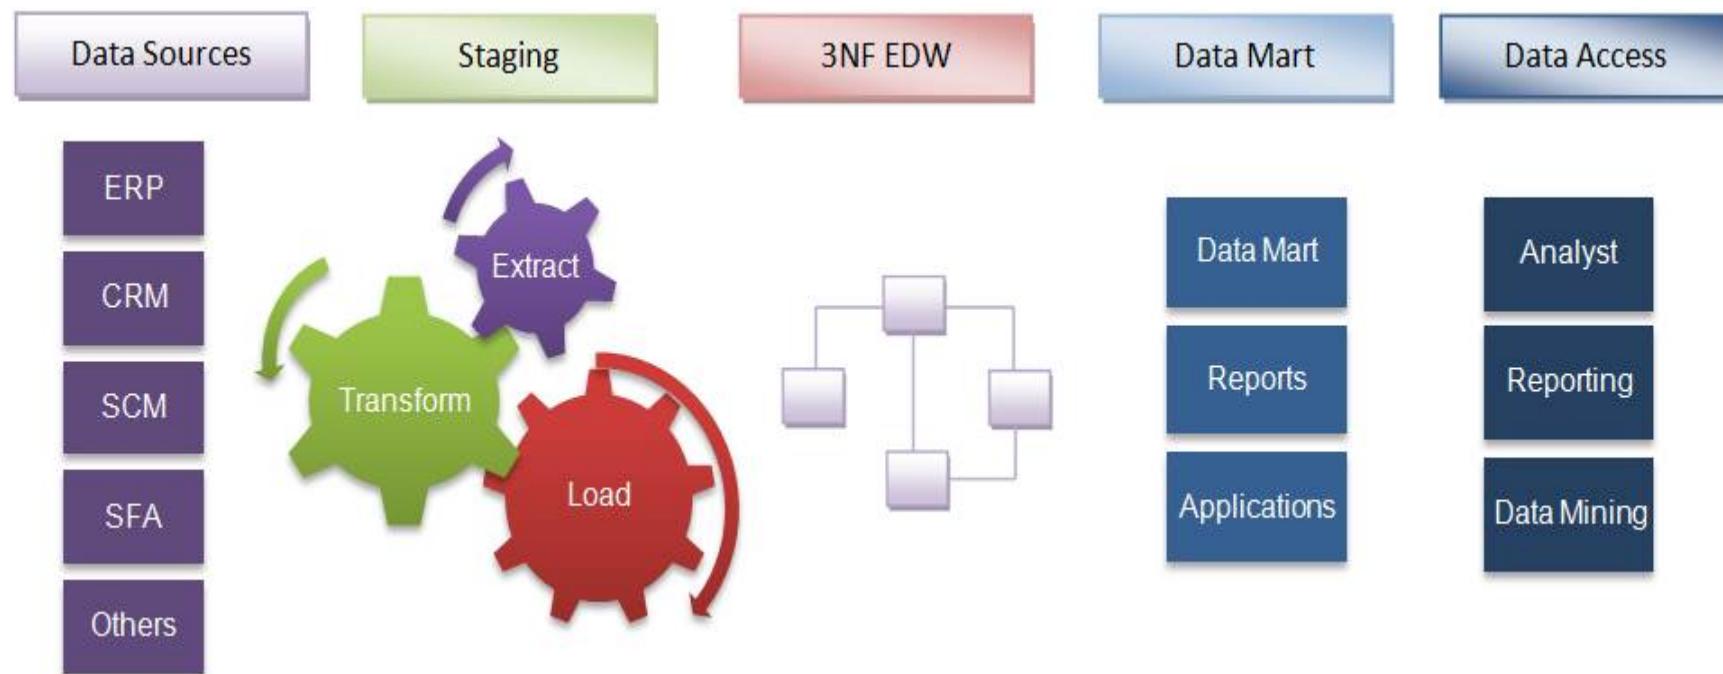
- **Bill Inmon's** enterprise data warehouse approach (the top-down design):

A normalized data model in 3NF is designed first. Then the dimensional data marts, which contain data required for specific business processes or specific departments are created from the data warehouse.

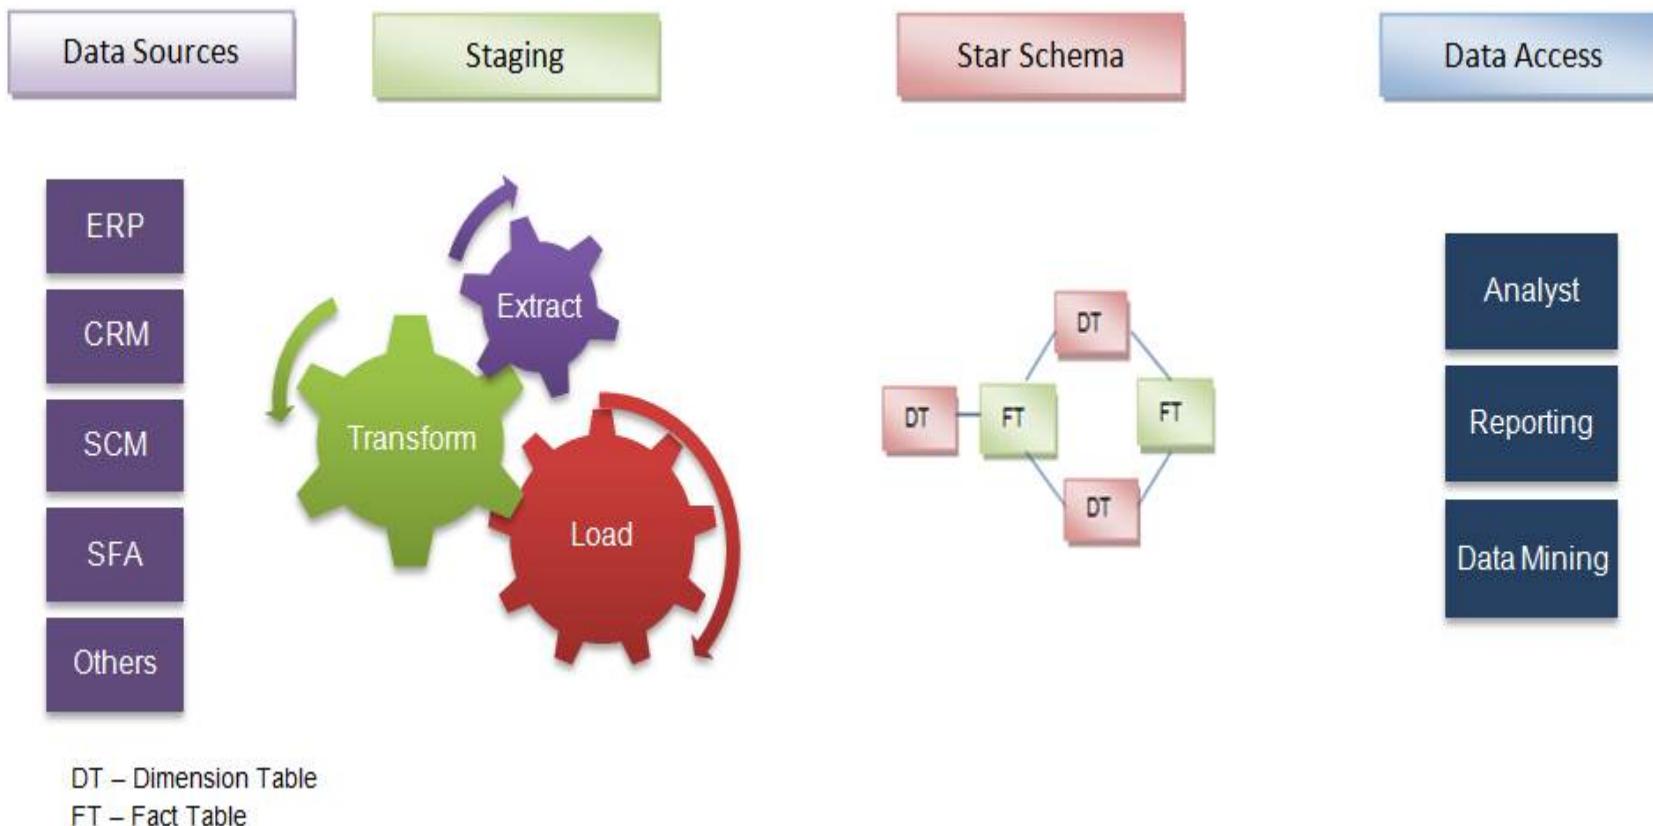
- **Ralph Kimball's** dimensional design approach (the bottom-up design):

The data marts for reporting and analysis are created first; these are then combined together to create a broad data warehouse.

Bill Inmon – Operational Data Store (ODS)

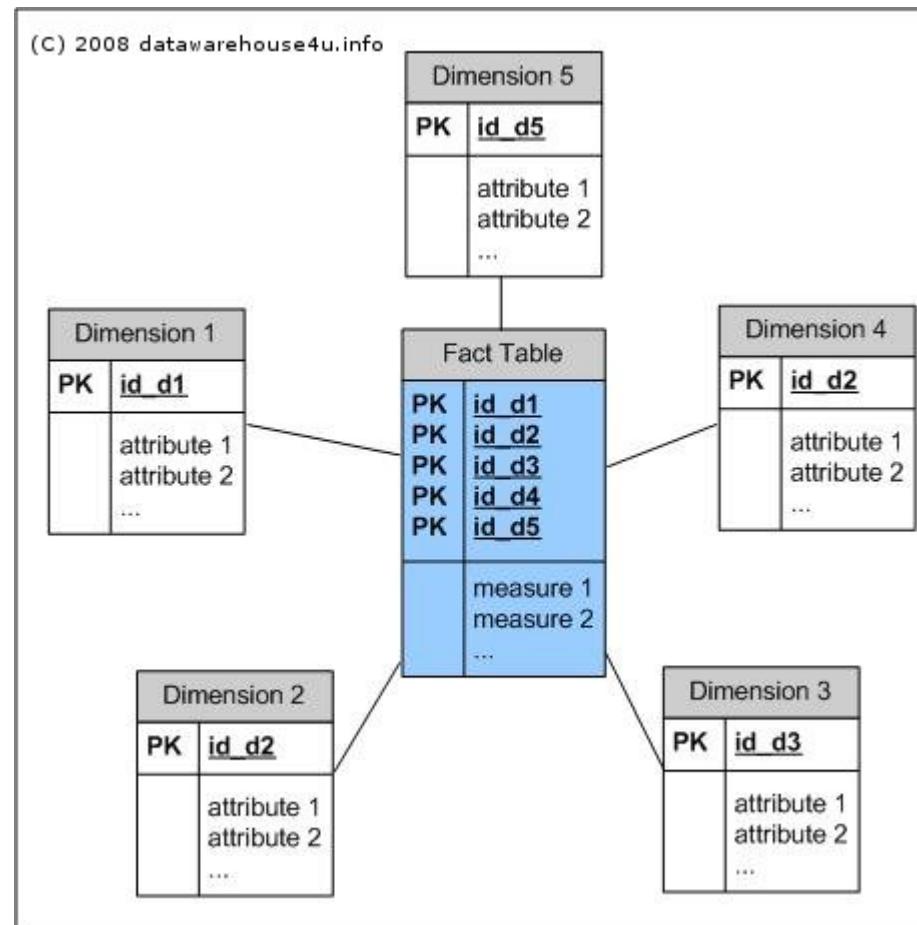


Ralph Kimball – Star Schema



Star Schema – Data Mart Design

Ralph Kimball



Star Schema – Data Mart Design

Ralph Kimball

The star schema architecture is the simplest [data](#) warehouse schema. It is called a star schema because the diagram resembles a star, with points radiating from a center.

The center of the star consists of fact [table](#) and the points of the star are the dimension tables.

Usually the fact tables in a star schema are in third normal form(3NF) whereas dimensional tables are de-normalized.

Despite the fact that the star schema is the simplest architecture, it is most commonly used nowadays and is recommended by [Oracle](#).

Pros and cons of both the approaches

	Inmon	Kimball
Building Data warehouse	Time Consuming	Takes lesser time
Maintenance	Easy	Difficult, often redundant and subject to revisions
Cost	High initial cost; Subsequent project development costs will be much lower	Low initial cost; Each subsequent phase will cost almost the same
Time	Longer start-up time	Shorter time for initial set-up
Skill Requirement	Specialist team	Generalist team
Data Integration requirements	Enterprise-wide	Individual business areas

Storing Data: NoSQL Databases

Unstructured Data Types

Here is a limited list of types of unstructured data:

- Emails
- Word Processing Files
- PDF files
- Spreadsheets
- Digital Images
- Video
- Audio
- Social Media Posts

NoSQL databases

NoSQL databases:

- Not using the relational model
- Schemaless
- Running well on clusters
- Tend to be open-source
- List of NoSQL databases (more than 150!):
<http://nosql-database.org/>

Why are NoSQL databases interesting?

Two primary reasons:

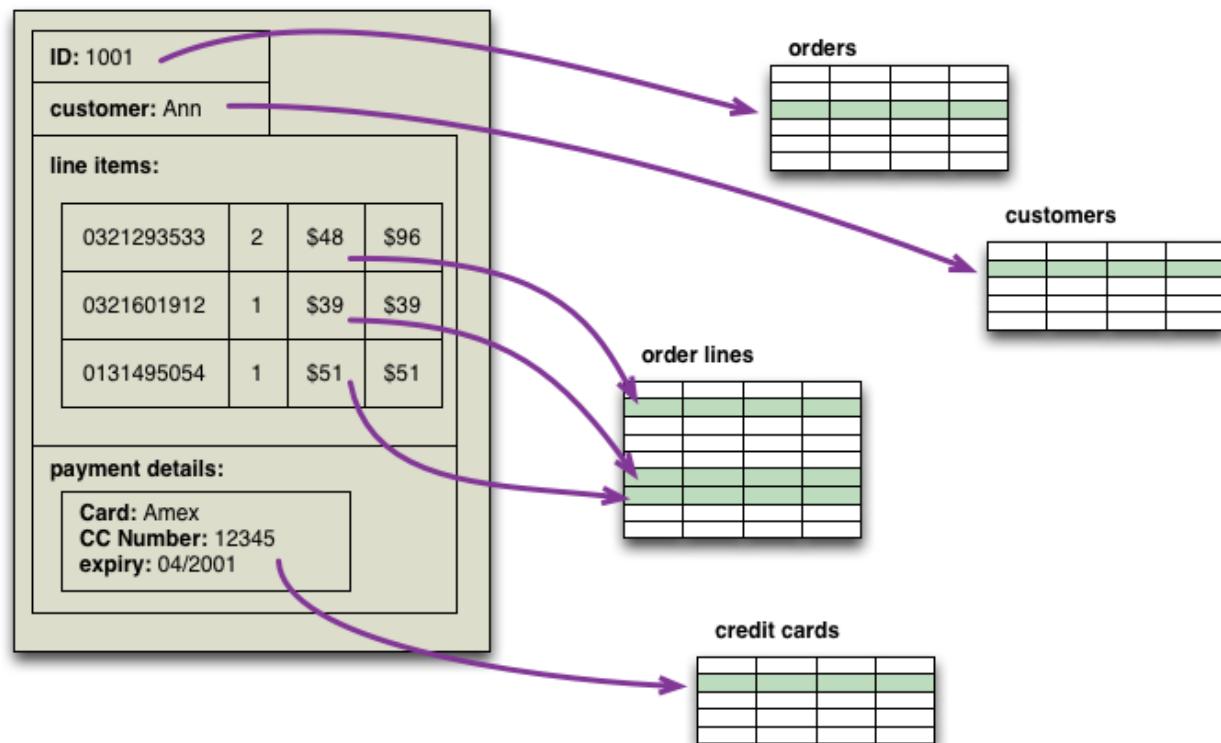
- Application development productivity
- Large-scale data

NoSQL databases: Data models

- **Key-value databases:**
 - BerkeleyDB, LevelDB, Memcached, Project Voldemort, Redis, Riak, ...
- **Document databases:**
 - CouchDb, MongoDB, OrientDB, RavenDB, Terrastore, ...
- **Column-family stores:**
 - Amazon SimpleDB, Cassandra, Hbase, Hypertable, ...
- **Graph databases:**
 - FlockDB, HyperGraphDB, Infinite Graph, Neo4J, OrientDB, ...
- **Other**

http://en.wikipedia.org/wiki/Data_warehouse

Aggregate data model: Example

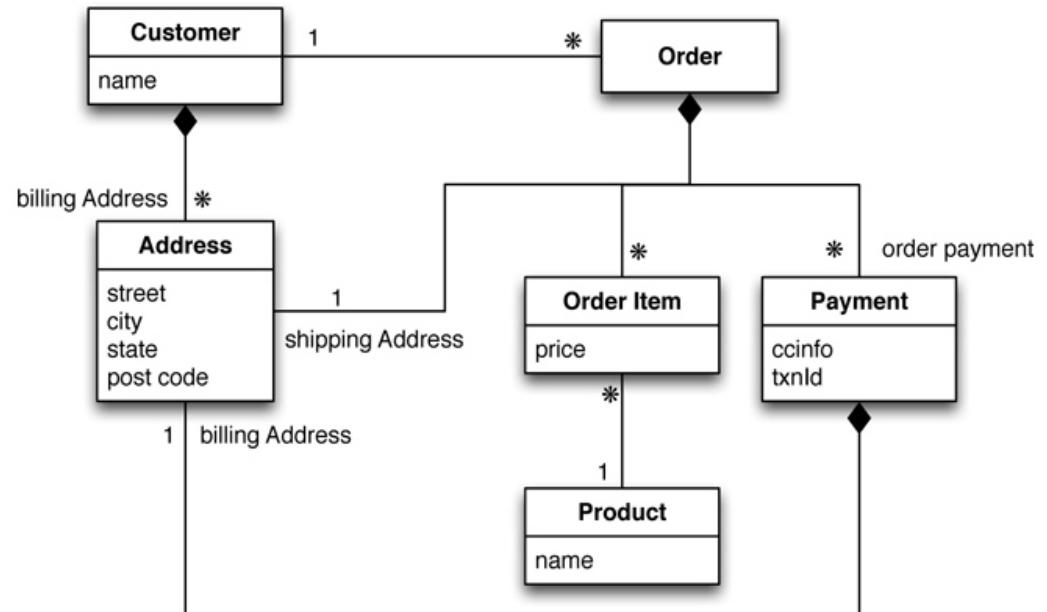


- **DDD: Domain-Driven Design**
- **Apparently easier to program**
- **Definitely, makes it easier to store and process data on multiple computer clusters.**

<http://martinfowler.com/bliki/AggregateOrientedDatabase.html>

Aggregate data model

```
// in customers
{
  "id":1,
  "name":"Martin",
  "billingAddress":[{"city":"Chicago"}]
}
// in orders
{
  "id":99,
  "customerId":1,
  "orderItems":[
    {
      "productId":27,
      "price": 32.45,
      "productName": "NoSQL Distilled"
    }
  ],
  "shippingAddress":[{"city":"Chicago"}]
  "orderPayment": [
    {
      "ccinfo":"1000-1000-1000-1000",
      "txnid":"abelif879rft",
      "billingAddress": {"city": "Chicago"}
    }
  ]
}
```

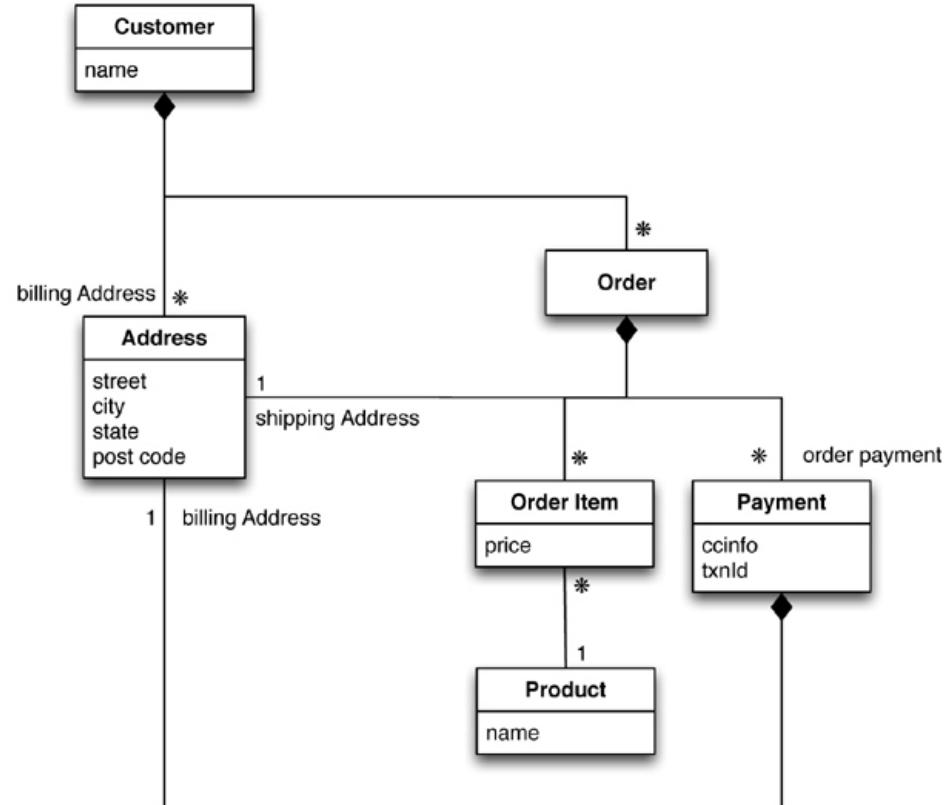


An aggregate data model

<http://martinfowler.com/bliki/AggregateOrientedDatabase.html>
 NoSQL distilled : a brief guide to the emerging world of polyglot persistence / Pramod J Sadalage, Martin Fowler.

Aggregate data model

```
// in customers
{
  "customer": {
    "id": 1,
    "name": "Martin",
    "billingAddress": [{"city": "Chicago"}],
    "orders": [
      {
        "id": 99,
        "customerId": 1,
        "orderItems": [
          {
            "productId": 27,
            "price": 32.45,
            "productName": "NoSQL Distilled"
          }
        ],
        "shippingAddress": [{"city": "Chicago"}]
      }
    ]
  }
}
```



An aggregate data model

<http://martinfowler.com/bliki/AggregateOrientedDatabase.html>
 NoSQL distilled : a brief guide to the emerging world of polyglot persistence / Pramod J Sadalage, Martin Fowler.

Aggregate data model: Pros and cons

- **Cons:**
 - It is often difficult to draw aggregate boundaries well
 - Does not support ACID transactions (thus sacrifices consistency)
 - Some queries are easier (to the point of being practical) but others may be really hard (e.g., to get to product sales history, you'll have to dig into every aggregate in the database).
- **Pros:**
 - Helps greatly with running on a cluster: This is the main argument for NoSQL databases.

<http://blog.dynatrace.com/2011/10/05/nosql-or-rdbms-are-we-asking-the-right-questions/>

Aggregate data model: Key points

- An aggregate is a collection of data that we interact with as a unit.
- Key-value, document, and column-family databases can all be seen as forms of aggregate-oriented database.
- Aggregates make it easier for the database to manage data storage over clusters.

Storing Data: NoSQL Database: Key-Value Databases

Key-value databases

- A key-value store is a simple hash table
 - Get the value for the key
 - Put a value for a key
 - Delete a key from the data store
- The value is a blob
- You can think of such databases as databases with only one table, which has two columns: ID and Value.

Key-value databases: Riak



Terminology comparison:

RDBMS	Riak
Database	Riak cluster
Table	Bucket
Row	Key-value
RowID (at least in Oracle)	key

1,000s of startups, enterprises, and organizations have deployed Riak for their production systems.



References: <http://wiki.basho.com/Riak.html>
NoSQL distilled : a brief guide to the emerging world of polyglot persistence / Pramod J Sadalage, Martin Fowler.



Data Storage

Key-value databases: Riak



- Writing to the Riak bucket using the *store* API:

```
Bucket bucket = getBucket(bucketName);
```

```
IRiakObject riakObject = bucket.store(key, value).execute();
```

- Getting value for the key using *fetch* API:

```
Bucket bucket = getBucket(bucketName);
```

```
IRiakObject riakObject = bucket.fetch(key).execute();
```

```
byte[] bytes = riakObject.getValue();
```

```
String value = new String(bytes);
```

References: <http://wiki.basho.com/Riak.html>

NoSQL distilled : a brief guide to the emerging world of polyglot persistence / Pramod J Sadalage, Martin Fowler.

Key-value databases: Riak



Riak provides an HTTP-based interface (this allows all operations to be performed from a web browser or command line):

- curl -X PUT HTTP://127.0.0.1:8098/riak/images/1.jpg -H "Content-type: image/jpeg" --data-binary @images.jpg
- curl -i HTTP://127.0.0.1:8098/riak/images/1.jpg
- curl -v -X POST -d '

```
{ "lastVisit":1324669989288,
"user":{"customerId":"1",
"name":"buyer",
"countryCode":"US",
"tzOffset":0}
}' -H "Content-Type:application/json" http://127.0.0.1:8098/riak/test/1
```
- curl -i HTTP://127.0.0.1:8098/riak/test/1

References: <http://wiki.basho.com/Riak.html>

NoSQL distilled : a brief guide to the emerging world of polyglot persistence / Pramod J Sadalage, Martin Fowler.

Key-value databases: Usage

- When to use:
 - Storing session information
 - User profiles, preferences
 - Shopping cart data
- When not to use
 - Relationships among data
 - Multi-operation transactions
 - Query by data
 - Operations by sets

References: <http://wiki.basho.com/Riak.html>
NoSQL distilled : a brief guide to the emerging world of polyglot persistence / Pramod J Sadalage, Martin Fowler.

Storing Data: NoSQL Database: Document Databases

Document databases

- **Documents are the main concept here**
 - DB stores and retrieves documents
- **Documents stored are similar to each other but do not have to be exactly the same**
 - Schemaless
- **Documents are stored in the value part of the key-value store**

Document databases: What is document?

- One document:
 - { "firstname": "Martin",
"likes": ["Biking", "Photography"],
"lastcity": "Boston",
"lastVisited":
}
- Another document:
 - {
"firstname": "Pramod",
"citiesvisited": ["Chicago",
"London", "Pune", "Bangalore"],
"addresses": [
{ "state": "AK",
"city": "DILLINGHAM",
"type": "R"
},
{ "state": "MH",
"city": "PUNE",
"type": "R" }
],
"lastcity": "Chicago"
}

Schema can differ significantly among documents in the same database.

This was not possible in RDBMS databases.

Document databases: MongoDB

Terminology comparison:



RDBMS	MongoDB
Database	MongoDB
Table	Collection
Row	Document
Rowid (at least Oracle)	_id
Join	DBRef



References: <http://www.mongodb.org/>
 NoSQL distilled : a brief guide to the emerging world of polyglot persistence / Pramod J Sadalage, Martin Fowler.

Document databases: MongoDB



- [DEMO] - <http://www.mongodb.org/#>
- Save document to MongoDB:
 - db.collectionName.insert({firstname: 'Martin', likes: ['Biking', 'Photography'], lastcity: 'Boston', lastVisited:"});
- Save another document to MongoDB:
 - db.collectionName.insert({firstname: 'Pramod', citiesvisited: ['Chicago', 'London', 'Pune', 'Bangalore'], addresses: [{ state: 'AK', city: 'DILLINGHAM', type: 'R'}, { state: 'MH', city: 'PUNE', type: 'R' }], lastcity: 'Chicago'});

References: <http://www.mongodb.org/>
NoSQL distilled : a brief guide to the emerging world of polyglot persistence / Pramod J Sadalage, Martin Fowler.

Document databases: MongoDB

[DEMO] - <http://www.mongodb.org/#>



Querying:

- All document in docName collection:

db.collectionName.find();

SQL: select * from docName

- Documents which satisfy a condition:

db.collectionName.find({firstname:"Martin"});

Equivalent to SQL query:

select * from collectionName where firstname = “Martin”

References: <http://www.mongodb.org/>
NoSQL distilled : a brief guide to the emerging world of
polyglot persistence / Pramod J Sadalage, Martin Fowler.

Document databases: MongoDB



- **SQL:**
 - **SELECT * FROM customerOrder, orderItem, product WHERE**
customerOrder.orderId = orderItem.customerOrderId
AND orderItem.productId = product.productId
AND product.name LIKE '%Big Data%'
- **MongoDB:**
 - **db.collectionName.find({"orders.productName":/Big Data/});**

References: <http://www.mongodb.org/>
NoSQL distilled : a brief guide to the emerging world of polyglot persistence / Pramod J Sadalage, Martin Fowler.

Document databases: Usage

- When to Use:
 - Event Logging
 - Content Management Systems, Blogging Platforms
 - Web Analytics or Real-Time Analytics
 - E-Commerce Applications
- When Not to Use
 - Complex Transactions Spanning Different Operations
 - Queries against Varying Aggregate Structure

Key-value vs. document data models

	Key-Value Model	Document Model
Aggregate	Opaque	Transparent
Access	Only lookup based on key (return whole aggregation)	Queries based on fields in the aggregate (can retrieve parts)
Index	-	Can create indexes based on the contents of the aggregate

<http://blog.dynatrace.com/2011/10/05/nosql-or-rdbms-are-we-asking-the-right-questions/>

Storing Data: NoSQL Database: Column-Family Stores

Column-oriented data model

Id	First Name	Last Name	Salary
1	Joe	Smith	40000
2	Mary	Jones	50000
3	Mike	Johnson	45000



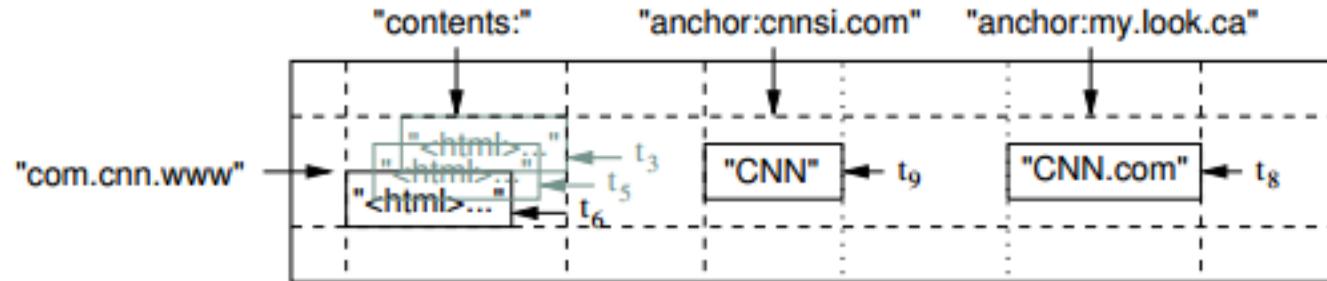
1, Joe, Smith, 40000
2, Mary, Jones, 50000
3, Mike, Johnson, 45000

1, 2, 3
Joe, Mary, Mike
40000,50000,45000

Column-oriented data model

- Column-oriented organizations are more efficient when an aggregate needs to be computed over many rows but only for a notably smaller subset of all columns of data, because reading that smaller subset of data can be faster than reading all data.
- Column-oriented organizations are more efficient when new values of a column are supplied for all rows at once, because that column data can be written efficiently and replace old column data without touching any other columns for the rows.
- Row-oriented organizations are more efficient when many columns of a single row are required at the same time, and when row-size is relatively small, as the entire row can be retrieved with a single disk seek.
- Row-oriented organizations are more efficient when writing a new row if all of the column data are supplied at the same time, as the entire row can be written with a single disk seek.

Column-oriented data model



- Data indexed by:
 - $(\text{row:string}, \text{column:string}, \text{time:int64}) \rightarrow \text{string}$
- #of distinct column families - small (in hundreds)
- unbounded number of columns
- Data processing is pushed to the application

Column-family databases

Conceptual View:

Row Key	Time Stamp	Column Family <i>contents</i>	Column Family <i>anchor</i>
“com.cnn.www”	t9		anchor:cnnsi.com=“CNN”
“com.cnn.www”	t8		anchor:my.look.ca=“CNN.com”
“com.cnn.www”	t6	contents:html=“<html>...”	
“com.cnn.www”	t5	contents:html=“<html>...”	
“com.cnn.www”	t3	contents:html=“<html>...”	

Physical View

Column Family *anchor*

Row Key	Time Stamp	Column Family <i>anchor</i>
“com.cnn.www”	t9	anchor:cnnsi.com=“CNN”
“com.cnn.www”	t8	anchor:my.look.ca=“CNN.com”

Column Family *contents*

Row Key	Time Stamp	Column Family <i>contents</i>
“com.cnn.www”	t6	contents:html=“<html>...”
“com.cnn.www”	t5	contents:html=“<html>...”
“com.cnn.www”	t3	contents:html=“<html>...”

<http://hbase.apache.org/book.html#datamodel>

Column-family databases: HBase

- **HBase is the Hadoop database**
 - Distributed, scalable, big data store
- Use HBase when you need random, real-time read/write access to your Big Data
- Goal is to host very large tables:
 - billions of rows X millions of columns
- HBase is an open-source, distributed, versioned, column-oriented store modeled after Google's Bigtable



Column-family databases: HBase

- No SQL-like query language
 - Java API
 - HBase Shell
 - » create ‘test’, ‘cf’
 - » put ‘test’, ‘row1’, ‘cf:a’, ‘value1’
 - » put ‘test’, ‘row2’, cf:b’, ‘value2’
 - » get ‘test’, ‘row1’
 - COLUMN CELL
 - cf:a timestamp=1288380727188,
value=value1
 - HBql – separate project to simplify the usage of Hbase
(hbql.com)



http://hbase.apache.org/book/quickstart.html#shell_exercises

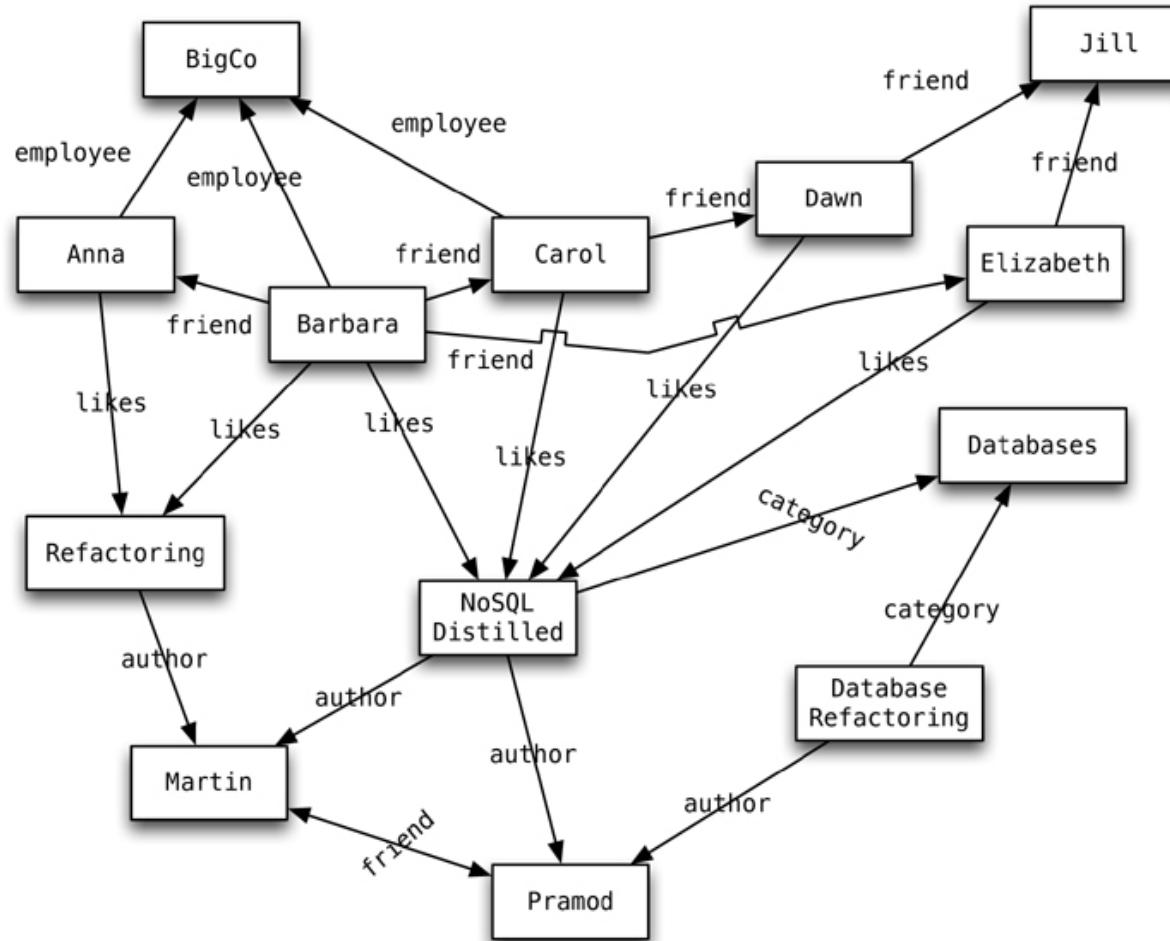
Column-family databases: Usage

- When to use:
 - Event logging
 - Content management systems, blogging platforms
 - Expiring usage
 - Need aggregate using, e.g., SUM or AVG
- When not to use
 - Frequent changes to the database (inserts and deletes maybe expensive)

References: <http://wiki.basho.com/Riak.html>;
NoSQL distilled : a brief guide to the emerging world of polyglot persistence / Pramod J Sadalage, Martin Fowler.

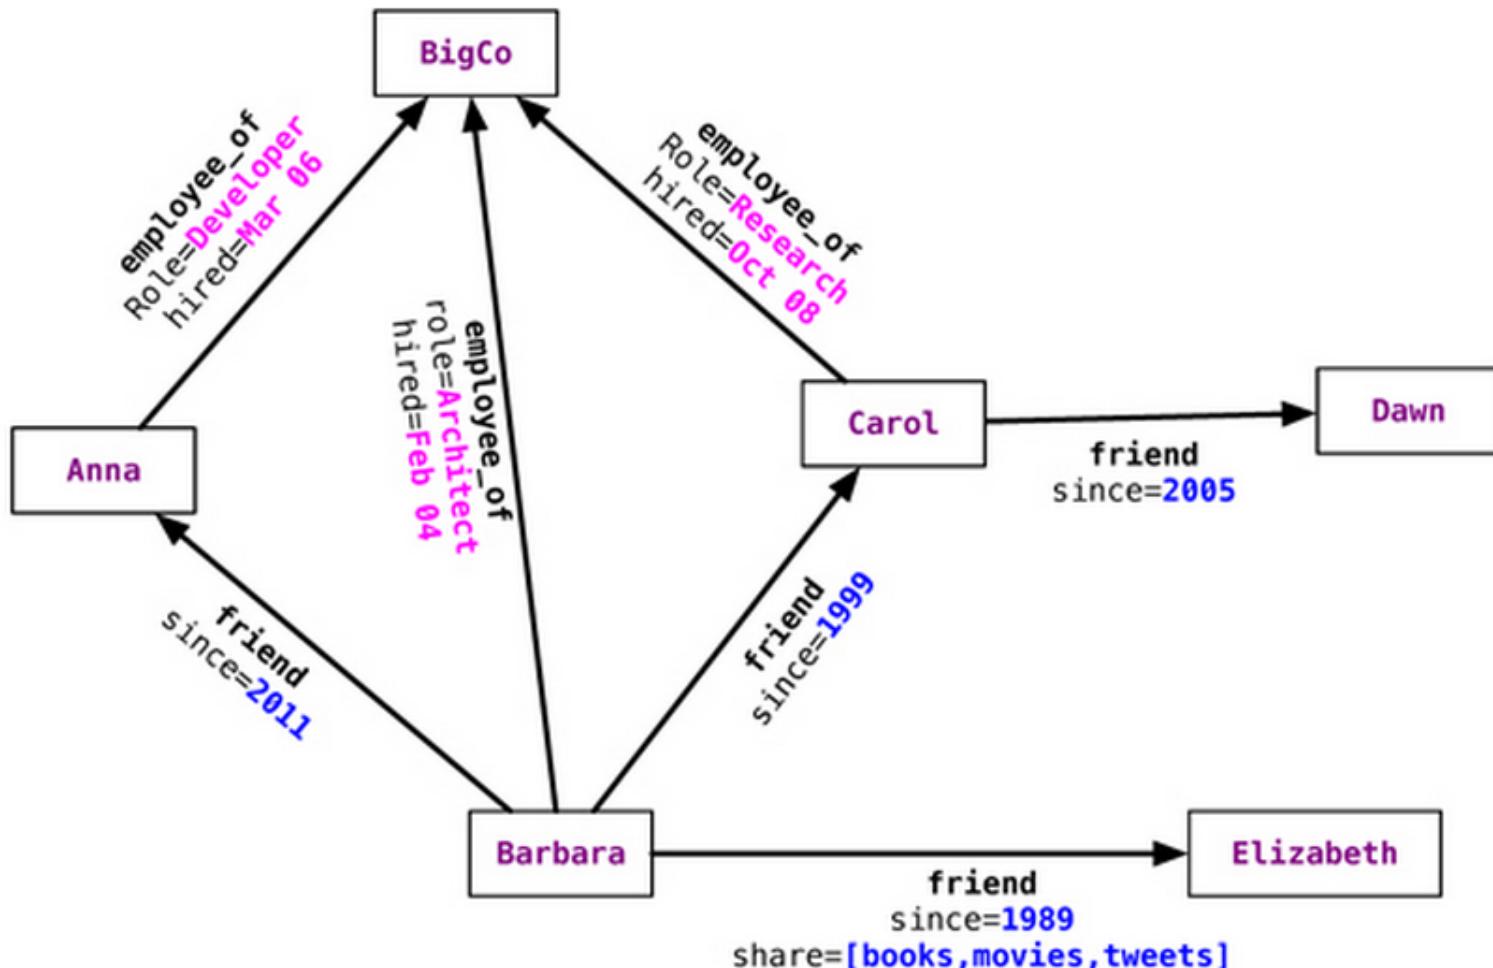
Storing Data: NoSQL Database: Graph Databases

Graph databases



An example graph structure

Graph databases



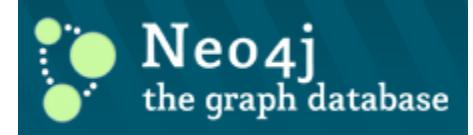
Relationships with properties

NoSQL distilled : a brief guide to the emerging world of polyglot persistence / Pramod J Sadalage, Martin Fowler.

Data Storage

Graph Databases: Neo4j

- Node creation:
 - `Node martin = graphDb.createNode();`
 - `martin.setProperty("name", "Martin");`
 - `Node pramod = graphDb.createNode();`
 - `pramod.setProperty("name", "Pramod");`
- Relationship creation:
 - `martin.createRelationshipTo(pramod, FRIEND);`
 - `pramod.createRelationshipTo(martin, FRIEND);`



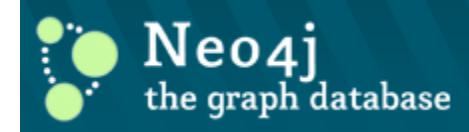
Graph Databases: Neo4j, Querying

Create index:

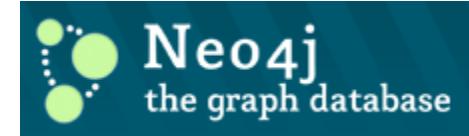
```
Transaction transaction = graphDb.beginTx();
try {
    Index<Node> nodeIndex = graphDb.index().forNodes("nodes");
    nodeIndex.add(martin, "name", martin.getProperty("name"));
    nodeIndex.add(pramod, "name", pramod.getProperty("name"));
    transaction.success();
} finally {
    transaction.finish();
}
```

Relationship creation:

```
Node martin = nodeIndex.get("name", "Martin").getSingle();
allRelationships = martin.getRelationships();
incomingRelations = martin.getRelationships(Direction.INCOMING);
```



Graph Databases: Neo4j, Querying



- Traverse the graphs at any depth:
 - `Node barbara = nodeIndex.get("name", "Barbara").getSingle();`
 - `Traverser friendsTraverser = barbara.traverse(Order.BREADTH_FIRST, StopEvaluator.END_OF_GRAPH, ReturnableEvaluator.ALL_BUT_START_NODE, EdgeType.FRIEND, Direction.OUTGOING);`
- Finding paths between two nodes:
 - `Node barbara = nodeIndex.get("name", "Barbara").getSingle();`
`Node jill = nodeIndex.get("name", "Jill").getSingle();`
`PathFinder<Path> finder = GraphAlgoFactory.allPaths(`
 `Traversal.expanderForTypes(FRIEND,Direction.OUTGOING)`
 `,MAX_DEPTH);`
`Iterable<Path> paths = finder.findAllPaths(barbara, jill);`

Graph Databases: Neo4j, Cypher QL

Cypher query language:

START beginingNode = (beginning node specification)

MATCH (relationship, pattern matches)

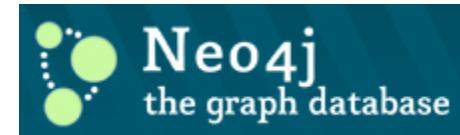
WHERE (filtering condition: on data in nodes and relationships)

RETURN (What to return: nodes, relationships, properties)

ORDER BY (properties to order by)

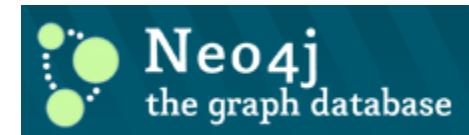
SKIP (nodes to skip from top)

LIMIT (limit results)



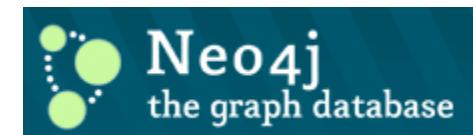
Graph Databases: Neo4j, Cypher QL

- Find all nodes connected to Barbara, either incoming or outgoing:
 - `START barbara = node:nodeIndex(name ="Barbara")
MATCH (barbara)--(connected_node)
RETURN connected_node`
- When we are interested in directional significance:
 - For incoming relationship
`MATCH (barbara) <-- (connected_node)`
 - For outgoing relationship
`MATCH (barbara) --> (connected_node)`



Graph Databases: Neo4j, Cypher QL

- Match can also be done on specific relationships using the `:RELATIONSHIP_TYPE` convention and returning the required fields or nodes:
 - `START barbara = node:nodeIndex(name = "Barbara")
MATCH (barbara)-[:FRIEND]->(friend_node)
RETURN friend_node.name,friend_node.location`
- Query for relationships where a particular relationship property exists:
 - `START barbara = node:nodeIndex(name = "Barbara")
MATCH (barbara)-[relation]->(related_node)
WHERE type(relation) = 'FRIEND'
RETURN related_node.name, relation.since`



Graph databases: Usage

- **When to Use:**
 - Connected Data
 - Routing, Dispatch, and Location-Based Services
 - Recommendation Engines
- **When Not to Use**
 - When you want to update all or a subset of entities
 - Not “graph” data model

NoSQL databases: Summary

- Key-value databases
- Document databases
- Column-family stores
- Graph databases

NoSQL Databases: Goals

- Not using the relational model
- Schemaless
- Running well on clusters
- Aggregates – nested data stored together

NoSQL: Schemaless

- Relational DB:
 - You first have to define a schema for your database
- NoSQL:
 - Storing data is more casual:
 - » Key-value store – allow any data under a key
 - » Document DB – no restrictions on the structure of the document
 - » Column-family DB – allow rows have different columns, any data under any column
 - » Graph DB – allow freely adding new edges and properties to nodes and edges

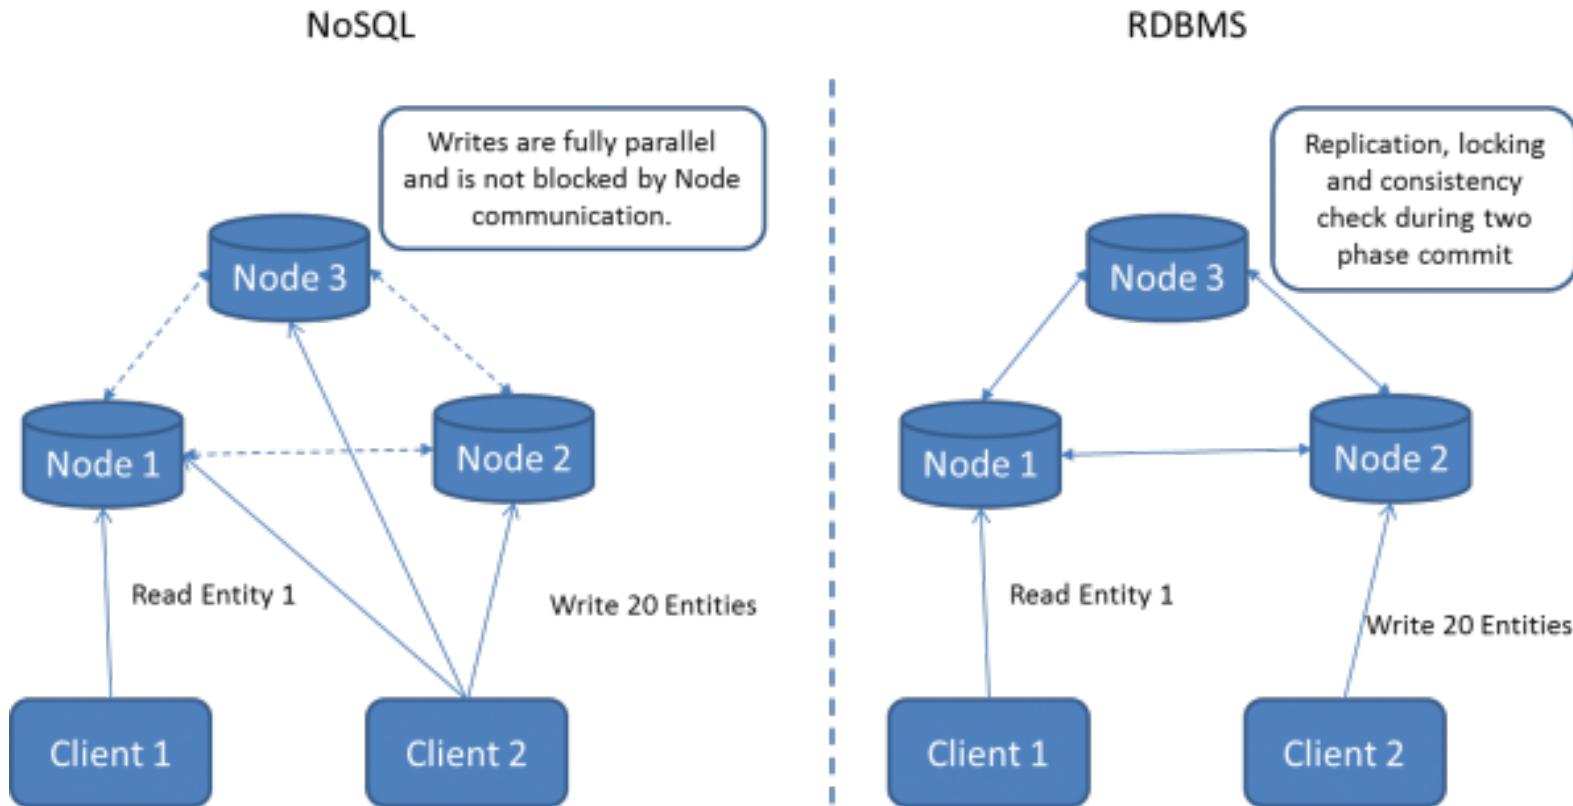
NoSQL: Schemaless

- **Advantages:**
 - Easy to handle changes
 - Easy to deal with non-uniform data
 - » where each record has different set of fields
- **Disadvantages:**
 - Database remains ignorant of the schema
 - » Can't validate data types
 - Implicit schema in the application code
 - » Bad and/or complicated code
 - » Multiple application access the same database

NoSQL Databases – BASE properties

- **Basically available:** Use replication to reduce the likelihood of data unavailability and use “sharding” (partitioning the data among many different storage servers) to make any remaining failures partial. The result is a system that is always available, even if subsets of the data become unavailable for short periods of time.
- **Soft state:** While ACID systems assume that data consistency is a hard requirement, NoSQL systems allow data to be inconsistent and relegate designing around such inconsistencies to application developers.
- **Eventually consistent:** Although applications must deal with instantaneous consistency, NoSQL systems ensure that at some future point in time the data assumes a consistent state. In contrast to ACID systems that enforce consistency at transaction commit, NoSQL guarantees consistency only at some undefined future time.

Why an RDBMS does not scale and many NoSQL solutions do?



NoSQL differs to RDBMS in the way entities get distributed and that no consistency is enforced across those entities

<http://blog.dynatrace.com/2011/10/05/nosql-or-rdbms-are-we-asking-the-right-questions/>

NoSQL databases

NoSQL databases:

- Not using the relational model ✓
- Schemaless ✓
- Running well on clusters ✓

Storing Data: Beyond NoSQL

Collecting data
Relational Databases
NoSQL Databases
● Beyond NoSQL Databases

Beyond NoSQL

- File systems
- XML Databases
- Object Databases
- Others ...

File systems

- Simple and widely implemented
 - Most of the devices have one or another file system
- More like key-value stores with hierachic key
 - No support for queries
- Little control over concurrency
 - Simple locking
- Cope with very large entities
 - Video, audio
- Very good for sequence access
- Works best for relatively small number of large files that can be processed in big chunks

XML databases

- Document-like databases
 - Documents are compatible with XML and various XML technologies are used to manipulate the document
- Allow to define schema
 - DTD, XML Schema
- Allow to perform transformation
 - XSLT
- Allow to query documents:
 - XPath, XQuery
 - SQL/XML

Object databases

- **Mapping from in-memory data structures to relational tables**
- **Close integration with the application**

Storage: Summary

When and why you (should) choose an RDBMS?

- Table based
- Relations between distinct Table Entities and Rows
- Referential Integrity
- ACID Transactions
- Arbitrary Queries and Joins

<http://blog.dynatrace.com/2011/10/05/nosql-or-rdbms-are-we-asking-the-right-questions/>

Storage: Summary

Why an RDBMS might not be right for you?

- If you just want to store your application entities in a persistent and consistent way
- If you have hierarchical application objects and need some query capability into them
- If you ever tried to store large trees or networks you will know that an RDBMS is not the best solution here
- If you are running in the Cloud and need to run a distributed database for durability and availability.
- You might already use a data warehouse for your analytics. If your data grows to large to be processed on a single machine, you might look into hadoop or any other solution that supports distributed Map/Reduce.

<http://blog.dynatrace.com/2011/10/05/nosql-or-rdbms-are-we-asking-the-right-questions/>

Storage: Summary Hybrid Systems? (examples)

F1 - The Fault-Tolerant Distributed RDBMS Supporting Google's Ad Business

Jeff Shute, Mircea Oancea, Stephan Ellner, Ben Handy, Eric Rollins, Bart Samwel, Radek Vingralek, Chad Whipkey, Xin Chen, Beat Jegerlehner, Kyle Littlefield, and Phoenix Tong. 2012. F1: the fault-tolerant distributed RDBMS supporting google's ad business. In Proceedings of the 2012 ACM SIGMOD International Conference on Management of Data (SIGMOD '12). ACM, New York, NY, USA, 777-778.
DOI=10.1145/2213836.2213954

<http://doi.acm.org/10.1145/2213836.2213954>

Oracle In-Database Hadoop: When MapReduce Meets RDBMS

Xueyuan Su and Garret Swart. 2012. Oracle in-database hadoop: when mapreduce meets RDBMS. In Proceedings of the 2012 ACM SIGMOD International Conference on Management of Data (SIGMOD '12). ACM, New York, NY, USA, 779-790.
DOI=10.1145/2213836.2213955

<http://doi.acm.org/10.1145/2213836.2213955>

<http://blog.dynatrace.com/2011/10/05/nosql-or-rdbms-are-we-asking-the-right-questions/>

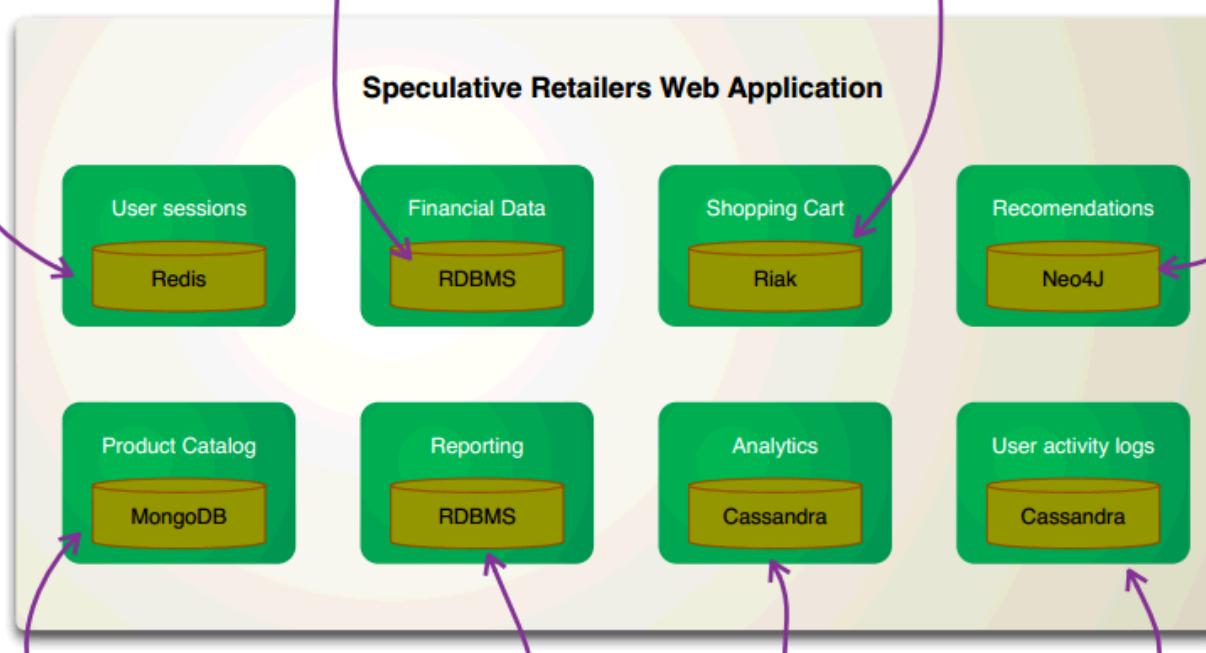
Polyglot Persistence

Rapid access for reads and writes. No need to be durable

Needs transactional updates. Tabular structure fits data

Needs high availability across multiple locations. Can merge inconsistent writes

Rapidly traverse links between friends, product purchases, and ratings



This is a very hypothetical example, we would not make technology recommendations without more contextual information

<http://martinfowler.com/articles/nosql-intro.pdf>

Suggested readings

Raghu Ramakrishnan, Johannes Gehrke. Database Management Systems. 3d Edition. WCB/McGraw-Hill 2003

Michael Stonebraker and Joseph M. Hellerstein , What Goes Around Comes Around (<http://mitpress.mit.edu/books/chapters/0262693143chapm1.pdf>)

NoSQL distilled: a brief guide to the emerging world of polyglot persistence / Pramod J Sadalage, Martin Fowler.

Tom White, Hadoop: The Definitive Guide, 3rd Edition, 2012

Fay Chang et. al, Bigtable: A Distributed Storage System for Structured Data

Jeffrey Dean and Sanjay Ghemawat. MapReduce: simplified data processing on large clusters. Commun. ACM 51, 1 (January 2008)

<http://hadoop.apache.org/>

http://hadoop.apache.org/docs/r0.20.2/hdfs_design.html

<http://hive.apache.org/docs/r0.9.0/>

<http://hbase.apache.org/>

<http://pig.apache.org/docs/r0.8.1/>

http://static.googleusercontent.com/external_content/untrusted_dlcp/research.google.com/en/us/pubs/archive/38125.pdf



DSL
Decision Systems Laboratory

Data Storage